


# MySQL Group Replication & MySQL InnoDB Cluster

Production Ready?

Kenny Gryp



productions

# Table of Contents

---

Group Replication

MySQL Shell (AdminAPI)

MySQL Group Replication

MySQL Router

Best Practices

Limitations

Production?

MySQL Group Replication

**MySQL InnoDB Cluster**

---

# MySQL Group Replication

---

- Developed by Oracle
- Generally Available in MySQL 5.7.17 on December 2016
- MySQL InnoDB Cluster as Solution

MySQL Group Replication is a MySQL Server plugin that **provides distributed state machine replication** with strong coordination between servers. Servers coordinate themselves automatically, when they are part of the same replication group. **Any server in the group can process updates. Conflicts are detected and handled automatically.** There is a **built-in membership service** that keeps the view of the group consistent and available for all servers at any given point in time. **Servers can leave and join the group** and the view will be updated accordingly.



# Asynchronous Replication vs. GR


---

## Async

- Async delivery
- Master -> Replica(s)
- Replica 'fetches' binlogs and executes
- external scripts required for automatic failover, split brain prevention...

## GR

- Sync delivery  
(at TRX Commit)
- Members <-> Members
- Majority of members receive TRX (PAXOS)
- Automatic handling of node status & membership, leader election (quorum-based)



# Group Replication

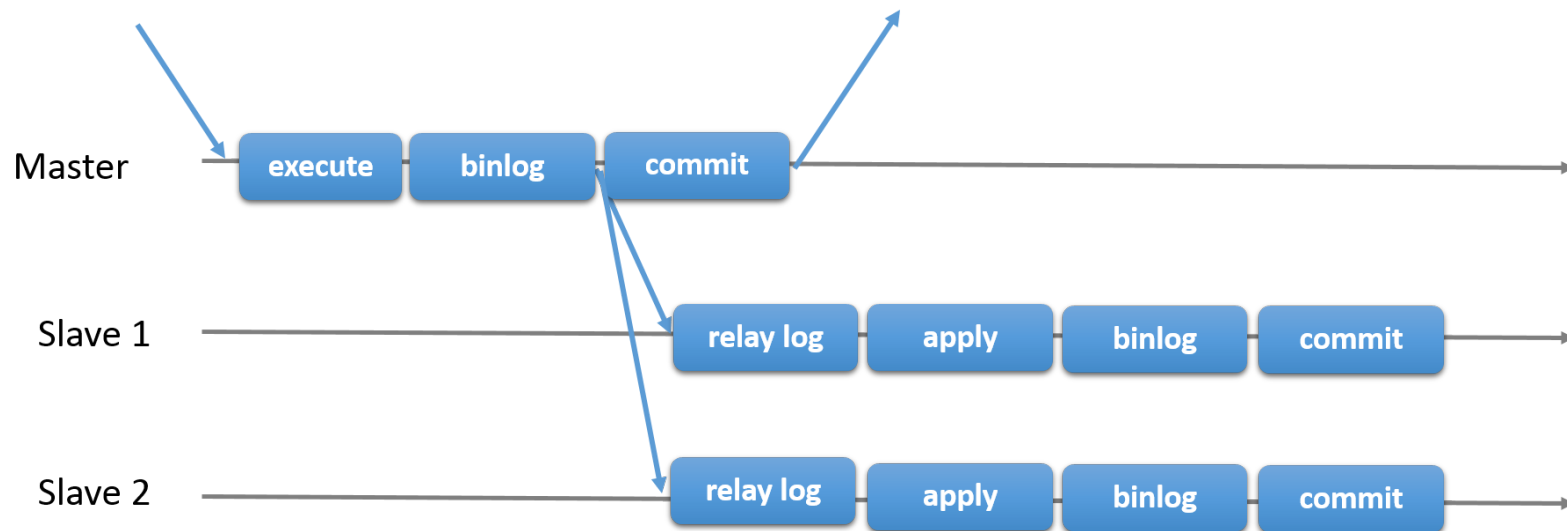
---

- GR uses a PAXOS protocol to ensure all nodes receive data
  - Increased COMMIT time  
similar to PXC (& semi-sync replication)
- Easy to configure/setup (easier than Async GTID Setups)
- (Integrated multi-node conflict detection)

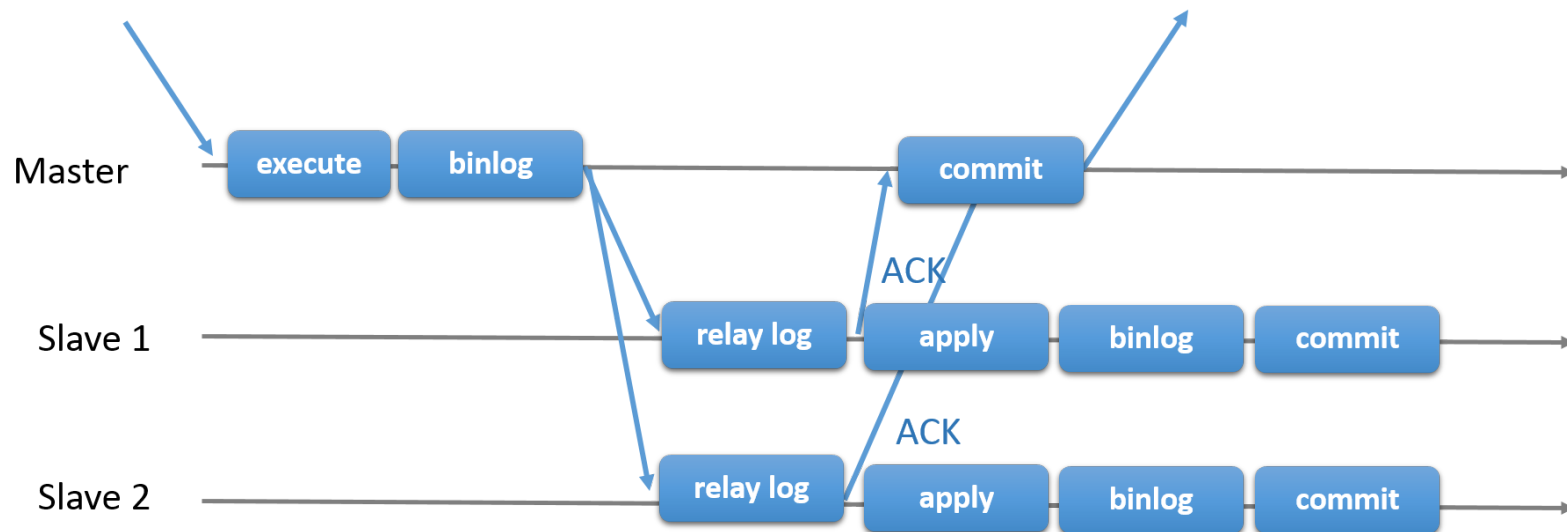


# Asynchronous Replication

---

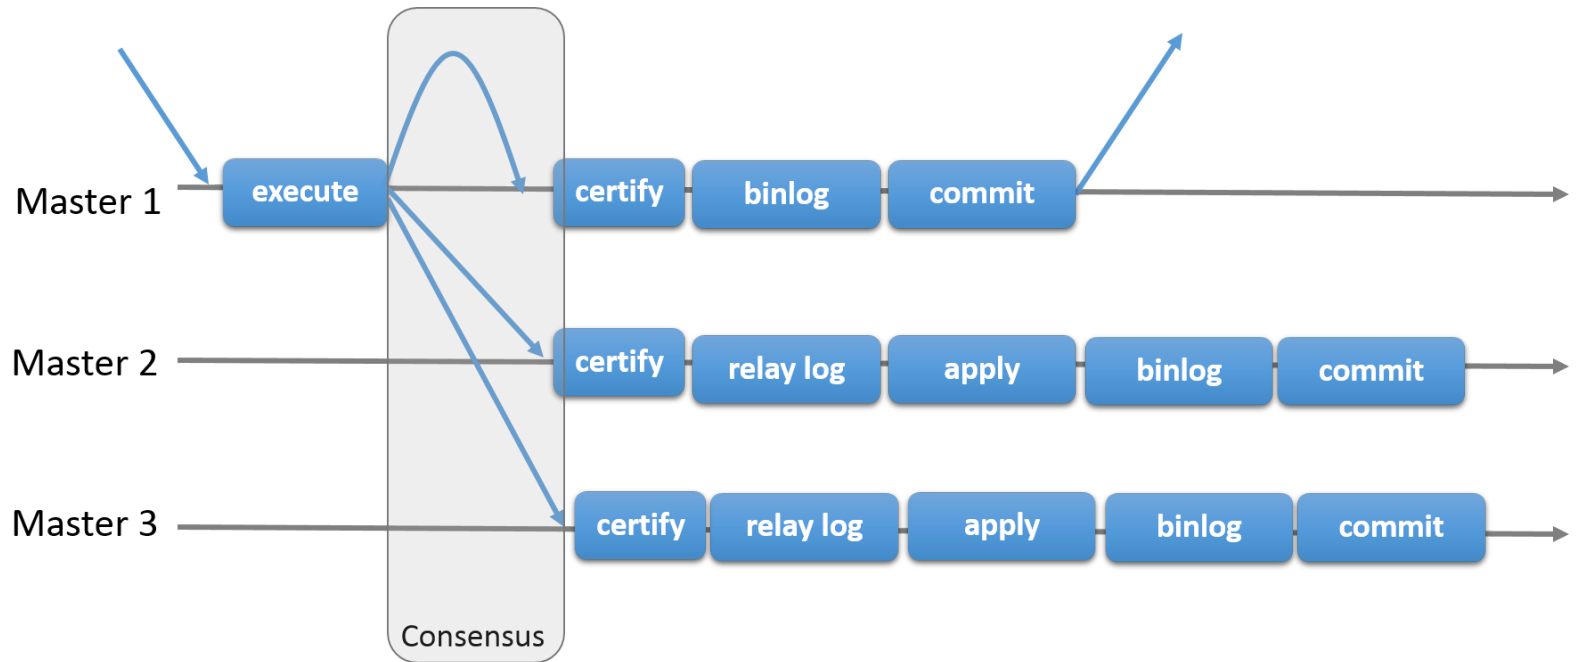


# Semi-Sync Replication





# Group Replication



# Use Cases

---


Environments Requiring:

- Strict **Durability** requirements
  - no data loss when a database node fails (0 RPO master failure):
  - **Consistency**: integrated split-brain prevention (Quorum based)
- Faster Failover than standard async (better RTO master failure)
- **(Write to multiple nodes simultaneously)**



# MySQL InnoDB Cluster

---



# MySQL InnoDB Cluster

---

Admin API

**MySQL Shell**

---

# MySQL Shell


---

```
mysql-js> var cluster  
          = dba.createCluster('prodCluster')
```

A new InnoDB cluster will be created  
on instance 'ic@ic-1:3306'.

Creating InnoDB cluster 'prodCluster' on 'ic@ic-1:3306'  
Adding Seed Instance...

Cluster successfully created. Use  
Cluster.addInstance() to add MySQL instances.  
At least 3 instances are needed for the cluster  
to be able to withstand up to one server failure.



# MySQL Shell

---

*"Makes Group Replication Configuration Easy"*


- Not really 5.7.21 & <= 8.0.4:
  - [#90439](#): AdminAPI does not change my.cnf
  - [#90438](#): AdminAPI fails to rejoin instances



# MySQL Shell

---

*"Makes Group Replication Configuration Easy"*

- Not really 5.7.21 & <= 8.0.4:
  - [#90439](#): AdminAPI does not change my.cnf
  - [#90438](#): AdminAPI fails to rejoin instances
- MySQL 8.0.11 (GA)
  -  (great unicode support)






# MySQL Shell

---

*"Makes Group Replication Configuration Easy"*

- Not really 5.7.21 & <= 8.0.4:
  - [#90439](#): AdminAPI does not change my.cnf
  - [#90438](#): AdminAPI fails to rejoin instances
- MySQL 8.0.11 (GA)
  -  (great unicode support)
  - Config is saved (SET PERSIST)
  - All actions can be done from a remote mysqlsh



# MySQL Group Replication

---

# MySQL Group Replication

---

- Split Brain Prevention
- Data Consistency
- Usability
- Stability
- Performance



Split Brain Prevention

# **MySQL Group Replication**

---

# Split Brain Prevention

---

No known split brain issues anymore!

Big improvement over 5.7.17  
(first GA)



Data Consistency

# **MySQL Group Replication**

---

# Data Consistency

---

Multi Writer

I have read the MySQL InnoDB cluster manual and  
I understand the requirements and limitations  
of advanced Multi-Master Mode.

Confirm [y/N]: NO



# Data Consistency

---

Multi Writer

I have read the MySQL InnoDB cluster manual and  
I understand the requirements and limitations  
of advanced Multi-Master Mode.

Confirm [y/N]: NO

Multi-Master is not recommended





# Data Consistency

---


## Multi Writer

I have read the MySQL InnoDB cluster manual and I understand the requirements and limitations of advanced Multi-Master Mode.

Confirm [y/N]: NO

Multi-Master is not recommended

- [#89194](#): Wrong certification lead to data inconsistency and GR breakage. (Multi-Master, should be fixed in 5.7.22 and 8.0.11)
- [#89938](#): Rejoin old primary node may duplicate key when recovery



Usability

# **MySQL Group Replication**

---

# Usability

---

```
mysql> INSERT INTO maurage  
      SELECT null FROM chez_lefred  
      WHERE dim0s_office IS NULL;  
ERROR 3100 (HY000): Error on observer while  
      running replication hook 'before_commit'.
```



# Usability

---

```
mysql> INSERT INTO maurage  
      SELECT null FROM chez_lefred  
      WHERE dim0s_office IS NULL;  
ERROR 3100 (HY000): Error on observer while  
      running replication hook 'before_commit'.
```


Error Log:

Plugin group\_replication reported:

'Error on session 75. Transaction of size 19943309 exceeds specified limit 15000000.'

To increase the limit please adjust group\_replication\_transaction\_size\_limit option.'

Run function 'before\_commit' in plugin 'group\_replication' failed



# Usability

---

```
mysql> COMMIT;
```

```
ERROR 1180 (HY000): Got error 149
```

```
- 'Lock deadlock; Retry transaction' during COMMIT
```



# Usability

---

```
mysql> COMMIT;
```

```
ERROR 1180 (HY000): Got error 149
```

```
- 'Lock deadlock; Retry transaction' during COMMIT
```

- Nothing in the error log!
- Cannot troubleshoot
- (Only happens in multi-writer mode)



# Usability

---

```
mysql> show processlist\G
```

Id: 25

User: root

Host: localhost

db: NULL

Command: Query

Time: 131

State: checking permissions

Info: create database node2



# Usability

---

```
mysql> show processlist\G
```

Id: 25

User: root

Host: localhost

db: NULL


Command: Query

Time: 131

State: checking permissions

Info: create database node2

- no Quorum
- `gr_unreachable_majority_timeout=0` by default  
:(





# Usability

---

## Features:

- No automatic node provisioning
- [#84730](#): Cannot troubleshoot Transaction Rollbacks
- [#90461](#): Changing replication mode cannot happen online
- [#84729](#): Impossible to block reads on partitioned nodes
- [#90484](#): No (easy) way to know if a GR node is writable or not
- [#90485](#): Ignore group\_replication\_group\_seeds nodes if they are not primary/active

## Bug:

- [#90483](#): DNS based seeds resolving to itself causes GR to not start



# Usability

---

Features & Bugs from Jean-François Gagné:

- [#89147](#): ... error messages is ambiguous.
- [#89145](#): Provide relay log details in case of Group Replication applier failure.
- [#89197](#): When GR fails, the error message says to "START SLAVE".



Stability

# **MySQL Group Replication**

---

# Stability

---

Feature:

- [#84784](#): Nodes do not reconnect back to the group replication once they got disconnected, causing nodes to drop from the cluster (except last 2 nodes)

Bug:

- [#90457](#): mysqld crash with ctrl-c/z'ed  
START GROUP\_REPLICATION



Performance

# **MySQL Group Replication**

---

# Performance

---

```
[ 220s] threads: 16 tps: 10599.99 qps: 10598.99 (r/w/o: 0.00/10598.99/0.00)
[ 221s] threads: 16 tps: 10571.71 qps: 10571.71 (r/w/o: 0.00/10571.71/0.00)
[ 222s] threads: 16 tps: 10307.88 qps: 10307.88 (r/w/o: 0.00/10307.88/0.00)
[ 223s] threads: 16 tps: 8220.26 qps: 8220.26 (r/w/o: 0.00/8220.26/0.00)
[ 224s] threads: 16 tps: 6381.09 qps: 6381.09 (r/w/o: 0.00/6381.09/0.00)
[ 225s] threads: 16 tps: 10348.85 qps: 10348.85 (r/w/o: 0.00/10348.85/0.00)
[ 226s] threads: 16 tps: 9383.95 qps: 9383.95 (r/w/o: 0.00/9383.95/0.00)
[ 227s] threads: 16 tps: 10528.06 qps: 10528.06 (r/w/o: 0.00/10528.06/0.00)

[ 280s] threads: 16 tps: 10335.09 qps: 10335.09 (r/w/o: 0.00/10335.09/0.00)
[ 281s] threads: 16 tps: 10372.06 qps: 10372.06 (r/w/o: 0.00/10372.06/0.00)
[ 282s] threads: 16 tps: 10237.61 qps: 10237.61 (r/w/o: 0.00/10237.61/0.00)
[ 283s] threads: 16 tps: 8206.20 qps: 8206.20 (r/w/o: 0.00/8206.20/0.00)
[ 284s] threads: 16 tps: 6050.79 qps: 6050.79 (r/w/o: 0.00/6050.79/0.00)
[ 285s] threads: 16 tps: 10053.31 qps: 10053.31 (r/w/o: 0.00/10053.31/0.00)
[ 286s] threads: 16 tps: 10208.14 qps: 10208.14 (r/w/o: 0.00/10208.14/0.00)
[ 287s] threads: 16 tps: 10315.78 qps: 10315.78 (r/w/o: 0.00/10315.78/0.00)
```



# Performance

---

```
[ 220s] threads: 16 tps: 10599.99 qps: 10598.99 (r/w/o: 0.00/10598.99/0.00)
[ 221s] threads: 16 tps: 10571.71 qps: 10571.71 (r/w/o: 0.00/10571.71/0.00)
[ 222s] threads: 16 tps: 10307.88 qps: 10307.88 (r/w/o: 0.00/10307.88/0.00)
[ 223s] threads: 16 tps: 8220.26 qps: 8220.26 (r/w/o: 0.00/8220.26/0.00)
[ 224s] threads: 16 tps: 6381.09 qps: 6381.09 (r/w/o: 0.00/6381.09/0.00)
[ 225s] threads: 16 tps: 10348.85 qps: 10348.85 (r/w/o: 0.00/10348.85/0.00)
[ 226s] threads: 16 tps: 9383.95 qps: 9383.95 (r/w/o: 0.00/9383.95/0.00)
[ 227s] threads: 16 tps: 10528.06 qps: 10528.06 (r/w/o: 0.00/10528.06/0.00)

[ 280s] threads: 16 tps: 10335.09 qps: 10335.09 (r/w/o: 0.00/10335.09/0.00)
[ 281s] threads: 16 tps: 10372.06 qps: 10372.06 (r/w/o: 0.00/10372.06/0.00)
[ 282s] threads: 16 tps: 10237.61 qps: 10237.61 (r/w/o: 0.00/10237.61/0.00)
[ 283s] threads: 16 tps: 8206.20 qps: 8206.20 (r/w/o: 0.00/8206.20/0.00)
[ 284s] threads: 16 tps: 6050.79 qps: 6050.79 (r/w/o: 0.00/6050.79/0.00)
[ 285s] threads: 16 tps: 10053.31 qps: 10053.31 (r/w/o: 0.00/10053.31/0.00)
[ 286s] threads: 16 tps: 10208.14 qps: 10208.14 (r/w/o: 0.00/10208.14/0.00)
[ 287s] threads: 16 tps: 10315.78 qps: 10315.78 (r/w/o: 0.00/10315.78/0.00)
```

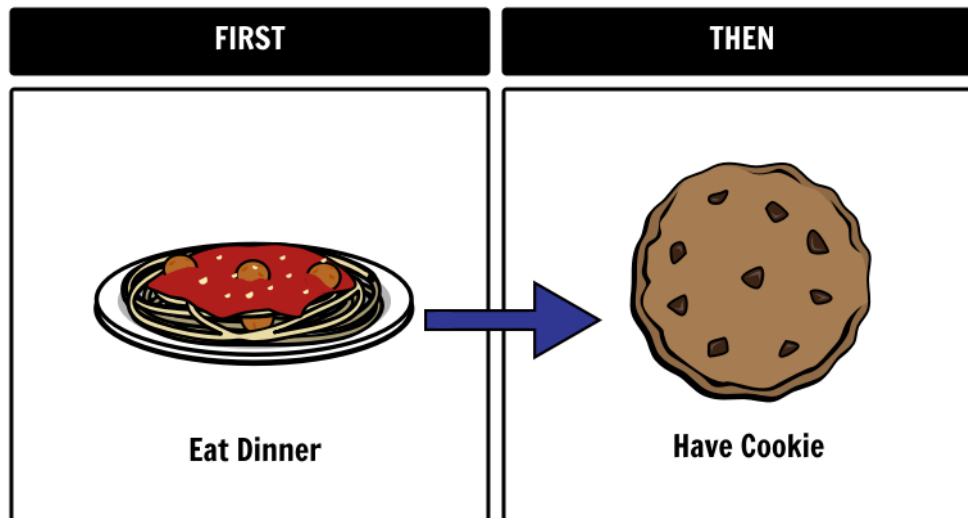
[#84774](#) Performance drop every 60 seconds




# Performance

---

Split-Brain Consistency & Usability first



Create your own at Storyboard That





# MySQL Router

---

# MySQL Router

---

- Quite simple load balancer:
  - TCP port for Writes & Reads
  - TCP port for Reads
- Routing Strategies (almost only valuable configuration setting)

first-available

next-available

round-robin

round-robin-with-fallback



# MySQL Router

---

- Quite simple load balancer:
  - TCP port for Writes & Reads
  - TCP port for Reads
- Routing Strategies (almost only valuable configuration setting)


first-available

next-available

round-robin

round-robin-with-fallback

- [#83236](#): Not possible to see mysqlrouter status  
*[quote]that's by design*  
*bugs.mysql.com is not a place to ask questions[/quote]*



# MySQL Router

---

## Limitations:

- No transparent read write splitting
- No query caching
- No connection multiplexing
- No way to get the router status
- No query rules
- No traffic mirroring
- No firewall



# MySQL Router

---

## Limitations:

- No transparent read write splitting
- No query caching
- No connection multiplexing
- No way to get the router status
- No query rules
- No traffic mirroring
- No firewall



Use ProxySQL!



# Best Practices

---

# Best Practices - Architecture

---

- Uneven amount of nodes
- Not recommended for WAN
  - => important timeouts are not configurable yet
- Use an intelligent Load Balancer
  - => [#84729](#) Impossible to block reads on partitioned nodes



# Best Practices - Configuration Settings

---

```
hostname=RESOLVABLE  
super_read_only=0N  
group_replication_unreachable_majority_timeout=20  
log_error_verbosity=3  
group_replication_ssl_mode=REQUIRED  
disabled_storage_engines="MyISAM, BLACKHOLE, FEDERATED,  
                           ARCHIVE, MEMORY"  
group_replication_auto_increment_increment=1
```






# Best Practices - Configuration Settings

---

```
hostname=RESOLVABLE
super_read_only=0N
group_replication_unreachable_majority_timeout=20
log_error_verbosity=3
group_replication_ssl_mode=REQUIRED
disabled_storage_engines="MyISAM, BLACKHOLE, FEDERATED,
                           ARCHIVE, MEMORY"
group_replication_auto_increment_increment=1
```

extra when using **5.7 & < 8.0.11**

```
group_replication_transaction_size_limit=150000000
group_replication_group_seeds=<ALL_NODES!>
group_replication_single_primary_mode=0N
group_replication_bootstrap_group=0FF
group_replication_allow_local_disjoint_gtids_join=0FF
```



# Best Practices

---

`hostname=VALID_RESOLVABLE_HOSTNAME`

other GR nodes will resolve the hostname to setup connections




# Best Practices

---

`super_read_only=0N`

Avoid PEBCAK split brain!

- Using `mysqlsh` with `< 8.0.11` does not persist configuration and GR does not start on boot
  - => writeable single mysql node when restarted




# Best Practices

---

`gr_unreachable_majority_timeout=20`

- Applications will get an error instead of hanging forever (Default 0)
- 20 seconds will abort group replication and configure `super_read_only=ON` (adapt to your needs)
- Drawback: if remaining 2 nodes get partitioned as well, all nodes go in ERROR and bootstrap is required



# Best Practices

---

`log_error_verbosity=3`

In MySQL 8, output is scarce, configure verbosity level 3 to allow better troubleshooting.



# Best Practices

---

`gr_ssl_mode=REQUIRED`

- DISABLED (default)
- Similar to client `ssl-mode=REQUIRED`
- `mysqlsh (py): dba.create_cluster('maurage',  
(memberSslMode='REQUIRED'))`



# Best Practices

---

```
disabled_storage_engines=  
"MyISAM, BLACKHOLE, FEDERATED,  
ARCHIVE, MEMORY"
```

Only InnoDB is supported!



# Best Practices

---

`gr_auto_increment_increment=1`

- Default 7
- Single-Primary/Writer is recommended
- No need for >1





# Best Practices

---

`gr_transaction_size_limit=150000000`

- < 8.0.2 default: unlimited maximum size of transactions
- >= 8.0.2 default: 143,051,147,460,9MB
- Keep Memory available for GR



# Best Practices

---

`gr_group_seeds=<ALL_NODES!>`

- `< 8.0.11`: with `mysqlsh` configured cluster does not properly configure seeds causing nodes not to rejoin [#90438](#)
- Configure IP Addresses, not hostnames [#90483](#)



# Best Practices

---

`gr_single_primary_mode=0N`

I have read the MySQL InnoDB cluster manual and  
I understand the requirements and limitations  
of advanced Multi-Master Mode.

Confirm [y/N]: NO




# Best Practices

---

`gr_bootstrap_group=OFF`

- Do not set this to ON, ONLY when creating a cluster.
  - does not go to OFF automatically
  - set back to OFF immediately
  - => use `dba.rebootClusterFromCompleteOutage('')` in some scenarios



# Best Practices

---

`gr_allow_local_disjoint_gtids_join=OFF`

- Don't even try to live with errant transactions
  - Big concern for data consistency
- Removed in 8.0.4



# Limitations

---

# Limitations

---

GR does not support:

- GET\_LOCK( )
- binlog\_format=STATEMENT
- Large transactions
- SELECT FOR UPDATE ([#85998](#))
- IPv6 ([#90217](#))
- Non InnoDB Storage Engines
- Consistent reads on all nodes
- No PK on all tables



**Production Ready?**

---



# Production Ready?

---

## Good

- Solid split brain prevention
- mysqlsh in 8.0.11 really starts to show it's power!



# Production Ready?

---

## Not So Good

- Many of the features listed in this presentation

## Bad

- [#84729](#): Impossible to block reads on partitioned nodes
- [#90484](#): No (easy) way to know if a GR node is writable
- Compared to Percona XtraDB Cluster/Galera Cluster:
  - No automatic node provisioning
  - Not possible to have synchronous reads

## Ugly

- [#84784](#): Nodes do not reconnect



# Production Ready? - My Opinion

---

(for the masses)

Component	MySQL 5.7 GA	MySQL 8.0 GA (+)
MySQL Shell	NO	YES
MySQL Router	NO (#)	NO (#)
Group Replication	NO (*)	NO (*)




# Production Ready? - My Opinion

---


(for the masses)

Component	MySQL 5.7 GA	MySQL 8.0 GA (+)
MySQL Shell	NO	YES
MySQL Router	NO (#)	NO (#)
Group Replication	NO (*)	NO (*)

(+) MySQL 8.0 is new, expect early adoption issues

(#)  Use ProxySQL!

(\*) **Early Adopters** required, much needed feedback to make the product better.




# Production Ready? - My Opinion

---

(for the masses)

Component	MySQL 5.7 GA	MySQL 8.0 GA (+)
MySQL Shell	NO	YES
MySQL Router	NO (#)	NO (#)
Group Replication	NO (*)	NO (*)

(+) MySQL 8.0 is new, expect early adoption issues

(#)  Use ProxySQL!

(\*) **Early Adopters** required, much needed feedback to make the product better.

- Best Practices!
- Test thoroughly!



# Ready For Production? (2018-04)

---



# Ready For Production? (2018-04)

---

- bottled end of 2016



# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful





# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful
- already very enjoyable for connoisseurs



# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful
- already very enjoyable for connoisseurs
- great legs



# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful
- already very enjoyable for connoisseurs
- great legs
- nice structure



# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful
- already very enjoyable for connoisseurs
- great legs
- nice structure
- needs some decanting to become top-knotch



# Ready For Production? (2018-04)

---



- bottled end of 2016
- delicious gem, still youthful
- already very enjoyable for connoisseurs
- great legs
- nice structure
- needs some decanting to become top-knotch
- KG: 90 points

